

Application: Model

Case study: Houses for sale

# Setup

In [9]:

```
%matplotlib inline

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.layers.experimental import preprocessing

print(tf.__version__)

sns.set_theme(style="ticks", color_codes=True)
```

2.4.1

## Data preparation

See notebook `10a-application-model-data-exploration.ipynb` for details about data preprocessing and data exploration.

In [10]:

```
ROOT = "https://raw.githubusercontent.com/kirenz/modern-statistics/main/data/"
DATA = "duke-forest.csv"
df = pd.read_csv(ROOT + DATA)

# Drop irrelevant features
df = df.drop(['url', 'address', 'type'], axis=1)

# Convert data types
df['heating'] = df['heating'].astype("category")
df['cooling'] = df['cooling'].astype("category")
df['parking'] = df['parking'].astype("category")

# drop column with too many missing values
df = df.drop(['hoa'], axis=1)

df.columns.tolist()
```

Out[10]:

```
['price',
 'bed',
 'bath',
 'area',
 'year_built',
 'heating',
 'cooling',
```

```
'parking',  
'lot']
```

# Simple regression

In [11]:

```
# Select features for simple regression
features = ['area']
X = df[features]

X.info()
print("Missing values:", X.isnull().any(axis = 1).sum())

# Create response
y = df["price"]
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 98 entries, 0 to 97

Data columns (total 1 columns):

#	Column	Non-Null Count	Dtype
0	area	98 non-null	int64

dtypes: int64 (1)

memory usage: 912.0 bytes

Missing values: 0

# Data splitting



In [12]:

```
from sklearn.model_selection import train_test_split

# Train Test Split
# Use random_state to make this notebook's output identical at every run
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

# Linear regression

Start with a single-variable linear regression, to predict price from area.

Training a model with `tf.keras` typically starts by defining the model architecture. In this case use a `keras.Sequential` model. This model represents a sequence of steps. In this case there is only one step:

- Apply a linear transformation to produce 1 output using `layers.Dense`.

The number of inputs can either be set by the `input_shape` argument, or automatically when the model is run for the first time.

In [13]:

```
lm = tf.keras.Sequential([
    layers.Dense(units=1, input_shape=(1,))
])

lm.summary()
```

Model: "sequential\_1"

---

Layer (type)	Output Shape
Param #	

dense_1 (Dense)	(None, 1)
2	

Total params: 2
Trainable params: 2

Non-trainable params: 0



This model will predict price from area.

Run the untrained model on the first 10 area values. The output won't be good, but you'll see that it has the expected shape, (10,1):

In [14]:

```
lm.predict(X_train[:10])
```

Out[14]:

```
array([[ -2667.3628],
       [ -2423.789  ],
       [-1526.7021],
       [-2527.6526],
       [-2145.2878],
       [-2129.6624],
       [-1997.3052],
       [-3606.7305],
       [-2290.513  ],
       [-2918.2898]], dtype=float32)
```

In [15]:

```
lm.compile(  
    optimizer=tf.optimizers.Adam(learning_rate=0.1),  
    loss='mean_absolute_error')
```

In [16]:

```
%%time
history = lm.fit(
    X_train, y_train,
    epochs=400,
    # suppress logging
    verbose=0,
    # Calculate validation results on 20% of the training data
    validation_split = 0.1)
```

CPU times: user 14.5 s, sys: 747 ms, total: 15.3 s

Wall time: 14.9 s



In [17]:

```
y_train
```

Out[17]:

49        525000

70        520000

68        412500

15        610000

39        535000

...

60        509620

71        540000

14        631500

92        590000

51        725000

Name: price, Length: 78, dtype: int64

In [18]:

```
# Calculate R squared
from sklearn.metrics import r2_score

y_pred = lm.predict(X_train).astype(np.int64)
y_true = y_train.astype(np.int64)

r2_score(y_train, y_pred)
```

Out[18]:

-0.8072585066424944

In [19]:

```
# slope coefficient  
lm.layers[0].kernel
```

Out[19]:

```
<tf.Variable 'dense_1/kernel:0' shape=(1, 1)  
dtype=float32, numpy=array([[118.29804]], dtype=  
float32)>
```

In [20]:

```
hist = pd.DataFrame(history.history)
hist['epoch'] = history.epoch
hist.tail()
```

Out[20]:


	loss	val_loss	epoch
395	248125.968750	194706.562500	395
396	247328.968750	193920.875000	396
397	246563.859375	193134.828125	397
398	245783.453125	192343.437500	398
399	244997.031250	191549.718750	399

In [21]:

```
def plot_loss(history):  
    plt.plot(history.history['loss'], label='loss')  
    plt.xlabel('Epoch')  
    plt.ylabel('Error [price]')  
    plt.legend()  
    plt.grid(True)
```

In [22]:

```
plot_loss(history)
```



Collect the results (mean squared error) on the test set, for later:

In [23]:

```
test_results = {}  
  
test_results['lm'] = lm.evaluate(  
    X_test,  
    y_test, verbose=0)  
  
test_results
```

Out[23]:

```
{ 'lm' : 276413.75 }
```

Since this is a single variable regression it's easy to look at the model's predictions as a function of the input:

In [26]:

```
x = tf.linspace(0.0, 6200, 6201)
y = lm.predict(x)

y
```

Out[26]:

```
array([[1.1863766e+02],
       [2.3693570e+02],
       [3.5523373e+02],
       ...,
       [7.3332988e+05],
       [7.3344819e+05],
       [7.3356650e+05]] , dtype=float32)
```




In [27]:

```
def plot_area(x, y):  
    plt.scatter(X_train['area'], y_train, label='Data')  
    plt.plot(x, y, color='k', label='Predictions')  
    plt.xlabel('area')  
    plt.ylabel('price')  
    plt.legend()
```

In [28]:

```
plot_area(x,y)
```



# Multiple Regression

In [29]:

```
# Select all relevant features
features= [
    'bed',
    'bath',
    'area',
    'year_built',
    'cooling',
    'lot'
]
X = df[features]

# Convert categorical to numeric
X = pd.get_dummies(X, columns=['cooling'], prefix='cooling', prefix_sep='_')

X.info()
print("Missing values:",X.isnull().any(axis = 1).sum())

# Create response
y = df["price"]
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 98 entries, 0 to 97

Data columns (total 7 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	bed	98 non-null	int64
1	bath	98 non-null	float64

2	area	98	non-null	int64
3	year_built	98	non-null	int64
4	lot	97	non-null	float64
5	cooling_central	98	non-null	uint8
6	cooling_other	98	non-null	uint8

dtypes: float64(2), int64(3), uint8(2)  
memory usage: 4.1 KB  
Missing values: 1

In [30]:

```
from sklearn.model_selection import train_test_split

# Train Test Split
# Use random_state to make this notebook's output identical at every run
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

In [31]:

```
lm_2 = tf.keras.Sequential([
    layers.Dense(units=1, input_shape=(7,))
])

lm_2.summary()
```

Model: "sequential\_2"

---

Layer (type)	Output Shape
Param #	

---

dense_2 (Dense)	(None, 1)
8	

---

---

Total params: 8  
Trainable params: 8

Non-trainable params: 0





In [32]:

```
lm_2.compile(  
    optimizer=tf.optimizers.Adam(learning_rate=0.1),  
    loss='mean_absolute_error')
```

In [33]:

```
%%time
history = lm_2.fit(
    X_train, y_train,
    epochs=400,
    # suppress logging
    verbose=0,
    # Calculate validation results on 20% of the training data
    validation_split = 0.1)
```

CPU times: user 14.5 s, sys: 739 ms, total: 15.2 s

Wall time: 14.7 s

In [34]:

```
# Calculate R squared
from sklearn.metrics import r2_score

y_pred = lm_2.predict(X_train).astype(np.int64)
y_true = y_train.astype(np.int64)

r2_score(y_train, y_pred)
```

Out[34]:

0.28772384550829944

In [35]:


```
# slope coefficients  
lm_2.layers[0].kernel
```

Out[35]:

```
<tf.Variable 'dense_2/kernel:0' shape=(7, 1)  
dtype=float32, numpy=  
array([[105.53154],  
       [108.65832],  
       [105.82798],  
       [104.10174],  
       [108.07097],  
       [104.49152],  
       [ 97.46925]], dtype=float32)>
```

In [36]:

```
plot_loss(history)
```



In [37]:

```
test_results['lm_2'] = lm_2.evaluate(  
    X_test, y_test, verbose=0)
```

DNN regression

The previous section implemented linear models for single and multiple inputs. This section implements a multiple-input DNN models. The code is basically the same except the model is expanded to include some "hidden" non-linear layers. The name "hidden" here just means not directly connected to the inputs or outputs.

These models will contain a few more layers than the linear model:

- Two hidden, nonlinear, Dense layers using the relu nonlinearity.
- A linear single-output layer.



In [38]:

```
dnn_model = keras.Sequential([
    layers.Dense(units=1, input_shape=(7,)),
    layers.Dense(64, activation='relu'),
    layers.Dense(64, activation='relu'),
    layers.Dense(1)
])

dnn_model.compile(loss='mean_absolute_error',
                  optimizer=tf.keras.optimizers.Adam(0.001))
```

In [39]:

```
%%time
history = dnn_model.fit(
    X_train, y_train,
    epochs=100,
    # suppress logging
    verbose=0,
    # Calculate validation results on 20% of the training data
    validation_split = 0.1)
```

CPU times: user 4.05 s, sys: 212 ms, total:  
4.26 s

Wall time: 4.09 s

In [40]:

```
# Calculate R squared
from sklearn.metrics import r2_score

y_pred = dnn_model.predict(X_train).astype(np.int64)
y_true = y_train.astype(np.int64)


r2_score(y_train, y_pred)
```

Out[40]:

0.4847519024408956

In [41]:

```
plot_loss(history)
```



In [42]:

```
test_results['dnn_model'] = dnn_model.evaluate(  
    X_test, y_test, verbose=0)
```

# Performance comparison

In [43]:

```
pd.DataFrame(test_results, index=['Mean absolute error [price]']).T
```

Out[43]:

Mean absolute error [price]	
lm	276413.750000
lm_2	157647.671875
dnn_model	143017.937500